

Social Science Methods for Legal Scholars

A Survey of Quantitative Text Analysis Methods

Joshua C. Fjelstul, Ph.D.

Post-Doctoral Research Fellow, University of Geneva

Researcher, ARENA Centre for European Studies, University of Oslo

What are we going to cover in this workshop?

- ▶ What is text analysis? What methods does it include?
- ▶ How does text analysis relate to statistics, machine learning (ML), and natural language processing (NLP)?
- ▶ How can we use text analysis to study international legal texts?
- ▶ How can we implement some basic text analysis tools in R?

What is quantitative text analysis?

- ▶ Quantitative text analysis is the use of quantitative methods to analyze the content of documents

Levels of analysis

- ▶ Corpus
- ▶ Documents
 - ▶ Books, documents, judgments, speeches, tweets
- ▶ Paragraphs
- ▶ Sentences
- ▶ Words

Assumptions

- ▶ The content of text reveals something meaningful and interesting about the author that we can use to answer research questions
- ▶ Text can be represented by features/tokens (words, lemmas, n -grams)
- ▶ The relative distribution of words in documents captures meaningful variation in topics and substantively important latent dimensions
 - ▶ We make a matrix called a document feature matrix (DFM) with words in columns, documents in rows, and frequencies in cells that describe these distributions
 - ▶ This might seem unsatisfactory
 - ▶ Some kinds of neural networks can take into account word order
 - ▶ Taking into account word order adds a lot of complexity and the "bag-of-words" approach usually works well

Types of quantitative text analysis

- ▶ Frequency analysis
 - ▶ What words/tokens are important?
 - ▶ Descriptive
- ▶ Similarity
 - ▶ How similar are documents/paragraphs/sentences?
 - ▶ Descriptive

Types of quantitative text analysis

► Scaling

- How do documents/paragraphs/sentences compare on a continuous latent dimension?
- Inferential (statistical models, machine learning models)

► Classification

- How do documents/paragraphs/sentences cluster into discrete groups?
- Inferential (statistical models, machine learning models)

Natural language processing (NLP)

- ▶ Text analysis overlaps with NLP — using computers to analyze language
 - ▶ Language processing
 - ▶ Optical character recognition (OCR)
 - ▶ Speech recognition
 - ▶ Morphology
 - ▶ Lemmatization
 - ▶ Stemming
 - ▶ Part-of-speech (POS) tagging
 - ▶ Semantics
 - ▶ Sentiment analysis
 - ▶ Named entity recognition
 - ▶ Word-sense disambiguation
 - ▶ High-level tasks
 - ▶ Machine translation
 - ▶ Natural language generation
 - ▶ Question answering

Machine learning (ML)

- ▶ Machine learning is the use of algorithms to predict outcomes based on data
- ▶ Algorithms learn the relationship between input data and labels based on training data and then make predictions for unseen data
- ▶ Quantitative text analysis is an application of ML to text data

- ▶ Supervised learning
 - ▶ Train a scaling/classification model on a pre-defined dimension using pre-coded training data
- ▶ Unsupervised learning
 - ▶ Train a scaling/classification model to endogenously learn an underlying dimension or categories using training data

- ▶ Supervised learning
 - ▶ Advantages
 - ▶ You already know the topics/dimensions so you know how to interpret your estimates
 - ▶ Disadvantages
 - ▶ It requires labeled training data (a lot of work)
 - ▶ You have to know ahead of time what the relevant topics/dimensions are

- ▶ Unsupervised learning

- ▶ Advantages

- ▶ You can explore naturally occurring topics and primary latent dimensions
 - ▶ You don't have to create labeled training data

- ▶ Disadvantages

- ▶ You have to validate that the topics/dimensions you uncover are meaningful and you have to figure out how to interpret them
 - ▶ You might not find meaningful topics/dimensions (or the ones you expect)
 - ▶ Unsupervised methods are harder to learn

Parametric vs non-parametric models

- ▶ Scaling and classification models can be parametric or non-parametric
- ▶ Parametric models overlap with statistics and involve estimating the values of parameters in statistical models that you can interpret to learn about the content of documents
- ▶ Non-parametric methods overlap with machine learning and depend on algorithms to learn the relationship between text data and labels

Classification

- ▶ Supervised methods for text
 - ▶ Naive Bayes classifier (non-parametric)
 - ▶ Regression classifiers (parametric)
 - ▶ Random forests (non-parametric)
 - ▶ Neural networks (non-parametric)
- ▶ Unsupervised methods for text
 - ▶ Latent Dirichlet allocation (LDA) (parametric)
 - ▶ Seeded LDA (parametric)
 - ▶ Structural topic models (STM) (parametric)

- ▶ Supervised
 - ▶ Wordscores (parametric)
- ▶ Unsupervised
 - ▶ Correspondence analysis (non-parametric) (dimensionality reduction)
 - ▶ Latent semantic analysis (non-parametric) (dimensionality reduction)
 - ▶ Wordfish (parametric)

Topic models

- ▶ Topics models are unsupervised methods for identifying naturally occurring topics in documents
- ▶ The most common type of topic model is latent Dirichlet allocation (LDA)
- ▶ It assumes that documents are mixtures of topics and that topics are mixtures of words
- ▶ Documents are not sorted into discrete categories
- ▶ A topic is a distribution of words
- ▶ The goal is to figure out which words are associated with which topics and which topics make up each document

Topic models

- ▶ The model estimates the probability that each topic applies to each document
- ▶ You have to tell the model how many topics to find
- ▶ You can look at the words that are most strongly associated with each topic
- ▶ Based on that list of words, we can label each topic
- ▶ Some topics will be more distinct than others
- ▶ If topics overlap too much, we may need fewer topics
- ▶ If topics are not distinct, we may need more topics

Word scores

- ▶ Supervised scaling method
- ▶ You start with a set of reference texts
- ▶ These need to represent the two poles of your latent dimension
- ▶ The reference texts are like a training set
- ▶ You calculate word scores based on the reference text and then use them to score the rest of the texts
- ▶ Each document will have a single score that represents its position on the latent dimension

Word scores

► Advantages

- After you choose the reference text, it's fully automated
- It scales all documents between the references texts at each end of the dimension

► Disadvantages

- It really matters which documents you use as the reference texts
- The dimension you define by choosing the reference texts may not be the dimension that explains the most variation in the content of your documents
- It's hard to choose the most extreme documents without a lot of knowledge about the content of documents (hard when there are a lot)

Wordfish (intuition)

- ▶ Unsupervised scaling method
- ▶ The input data is a DFM
- ▶ We don't have to know the underlying dimension ahead of time
- ▶ So we have to show that our estimates capture a meaningful latent dimension
- ▶ Based on a poisson distribution
 - ▶ The poisson distribution models counts of discrete events — like the occurrence of words in a document
 - ▶ Wordfish is a type of poisson scaling model

Wordfish (equation)

- ▶ The model equation is:

$$\log \lambda_{ij} = \alpha_i + \theta_i \beta_j + \phi_j$$

- ▶ Parameters:

- ▶ λ_{ij} is the expected frequency of a word i in document j
- ▶ i indexes documents
- ▶ j indexes words
- ▶ θ_i is the latent position of document i that we want to estimate
- ▶ β_j is the latent position of word j (strength of the relationship between the latent dimension and the frequency of the word)
- ▶ α_i is a document fixed effect (controls for how long each document is)
- ▶ ϕ_j is a word fixed effect (controls for how frequent each word is)

Wordfish (estimation)

- ▶ On the right-hand side of a regression equation, we have data and parameters
- ▶ The data is constant and we estimate the parameters
- ▶ But here, there's no data on the right-hand side, so what do we do?
 - ▶ We start with random values for all of the parameters
 - ▶ We hold ϕ and β (the word parameters) constant and estimate α and θ (the document parameters)
 - ▶ Then we hold α and θ (the document parameters) constant and estimate ϕ and β (the word parameters)
 - ▶ And we iterate back and forth
 - ▶ This is called expectation-maximization
 - ▶ Eventually, we'll converge to good estimates of all parameters

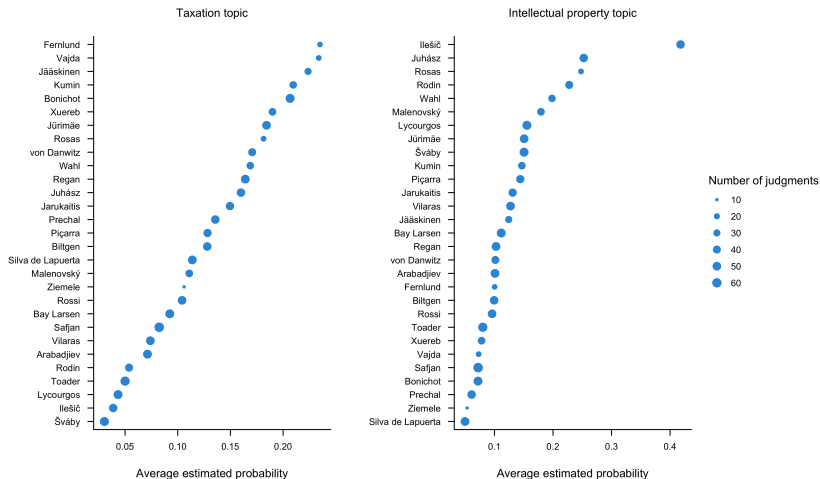
Wordfish (interpretation)

- ▶ We can interpret the θ estimates as the position of each document on our latent dimension. This is the main thing we're interested in
- ▶ But this is an unsupervised model, so how do we know what the latent dimension is?
 - ▶ We can plot the word fixed effects (y-axis) against the word positions (x-axis) to get an idea of what the dimension is
 - ▶ Rarer words (lower fixed effect) will be more discriminatory (will provide more information about how documents differ)
 - ▶ Words with more extreme positions (x-axis) will define the substantive content of each pole of the latent dimension
 - ▶ You have to convince your audience that your latent dimension is meaningful and interesting


Research questions

- ▶ We'll answer two research questions about the content of judgments delivered by the Court of Justice of the European Union (CJEU)
 - ▶ To what extent do judges specialize in certain areas of law?
 - ▶ What is the primary latent dimension in CJEU judgments? Is it a left/right dimension? Is it a pro-/anti-European integration dimension? Or is it something else, like a policy/procedure dimension?
- ▶ We'll use unsupervised and semi-supervised topic models to address the first question and an unsupervised scaling model to address the second question

Seeded LDA model



Wordfish (word positions)



Wordfish (document positions)

